

# The Blockchain Backbone Model

Presented by Dionysis Zindros

Invented by **Juan Garay, Aggelos Kiayias, Nikos Leonardos**  
Alcoholics - National Technical University of Athens, November 2018

# What we'll talk about

- The **blockchain backbone**
- The **mathematical theory** that governs **bitcoin**,  
ethereum, litecoin, monero, bitcoin cash, ethereum classic, zcash, doge  
and hundreds of other cryptocurrencies  
(Originally called the “*bitcoin* backbone”, but really quite generic)
- The **algorithmic parts** of blockchains
- The **consensus layer** of blockchains
- The nature of a blockchain: What a blockchain is
- What **problem** do blockchains solve?
- **Proof-of-work** blockchains only!  
(Not proof-of-stake such as Cardano, Decred)


# What we **won't** talk about

- Blockchains are a complicated subject
- Many engineering and science problems are open
- We won't talk about **engineering challenges**
- We will only speak about **one model**
- Here we treat the problem in the closed setting of **computer science theory**
- We won't be concerned with **implementation details**
- Blockchains deal with a lot more than the **consensus layer**, including the **application layer**
- We won't speak about **specific cryptocurrencies**
- We will make many **simplifying assumptions**
- We won't **prove** anything today -- I will talk only about definitions


# The Blockchain Backbone

- The *first theory* formalizing the security of bitcoin & cryptocurrencies
- Best model we have so far
- It answers the question: **Why is bitcoin secure?**
- This question was not answered in the original Bitcoin paper!
- Invented and published in EUROCRYPT '15 by Juan Garay, Aggelos Kiayias, and Nikos Leonardos
- Revisited and published in CRYPTO '17 by same authors
- One of the most significant works after Nakamoto's original paper

# A peer-to-peer network of money



# A peer-to-peer network of money



# What's a peer-to-peer network of money?


- Every node on the network knows how much money is in **every account**
- When someone pays someone else, they generate a *transaction* -- **tx**
- The tx subtracts money from one account and adds it to another
- tx are *broadcast to everyone*
- By receiving tx, nodes can update their view of others' accounts
- A tx can only be created by the owner sending the money:  
It needs a digital signature

Most importantly: The network is **decentralized**. There is **no trusted third party!**  
(Everything is very easy if we have a trusted bank or PayPal)

# Assumption: Time is discrete

We think of time as happening in **discrete rounds**: 1, 2, 3, ...,  $r$

Information (transactions) appears in rounds



The network is **synchronous**:


A message sent by an honest party during round  $r$  is received at round  $r + 1$  by all other honest parties. One round in practice corresponds to the time needed to traverse the whole network.




# The adversary

- The adversary can be **anything**
- Modelled as an arbitrary probabilistic polynomial-time Turing Machine
- This analysis is powerful:  
We do not say what attacks we have in mind!
- The protocol is secure against **any attack** we may not even think of!
- There is only **one adversary**
- The adversary **controls** multiple parties
- All of them coordinate together to attack the protocol
- Parties controlled by the adversary can communicate arbitrarily *during* a round

# Assumption: The adversary controls arbitrary nodes



# Assumption: The network is Sybil attackable



# The Sybil attack

The Sybil attack captures the fact that the adversary can:

- Fake IP addresses
- Fake e-mail addresses, Facebook accounts, Gmail accounts
- Fake phone numbers

Generally, we cannot rely on any centralized service or network heuristic!  
We cannot simply take a vote by counting.

# The parties

- There are  $n$  parties in total
- $t < n$  are controlled by **one** adversary (corrupted)
- $n - t < n$  are honest
- Each party wishes to publish some  $tx$  to the rest
- There's new tx coming in by honest parties during various rounds

# The consensus problem

- It may take some time for tx to travel on the network
- How can the parties agree **which tx happened first?**
- This question is crucial for security:  
If a tx did/didn't happen, a party may/may not have sufficient money!
- **tx validity depends on past tx**

time  $t = 1$



**Eve** really has 1 BTC



**Bob** & **Alice** think  
Eve has 1 BTC

time  $t = 2$



**Eve** spends her 1 BTC  
to buy coffee from Alice

Sends transaction  $tx_1$  to **Alice**



...simultaneously...


time  $t = 2$



**Eve** spends the *same* 1 BTC  
to pay back **herself**

Sends transaction  $tx_2$  to **Bob**





**Alice** and **Bob** exchange the transactions on the network  
and see a **double spend!**

The two transactions are irreconcilable.  
Which one is the right one?



# Simple ideas don't work

**“Oldest transaction is the right one.  
Reject the newer double spend!”**

**Bob** and **Alice** don't agree on which transaction happened first!  
**Alice** thinks  $tx_1$  happened first. **Bob** thinks  $tx_2$  happened first.

# Simple ideas don't work

**“If you see a double spend, you know it's a bad actor!  
Reject the transaction!”**

How long do you wait until you know there has been *no double spend*?

What if **Eve** bought a cup of coffee from **Alice**?

When does **Alice** deliver the coffee?


# Simple ideas don't work

“Wait for some constant time  $\Delta$ .

If no double spend has occurred within time  $\Delta$ , accept the transaction.

If a double spend occurs within time  $\Delta$ , reject both transactions.

If a double spend occurs after time  $\Delta$ , reject the double spend.”



**Eve** creates tx<sub>1</sub> and sends it to both **Alice** and **Bob** at time  $t = 0$ .

Then **Eve** sends tx<sub>2</sub> to Bob at time  $t = \Delta - 1$ , rendering tx<sub>1</sub> invalid according to **Bob**.

Eve sends tx<sub>2</sub> to **Alice** at time  $t = \Delta + 1$ , rendering tx<sub>1</sub> valid according to **Alice**.

**Alice** thinks tx<sub>1</sub> is valid, but **Bob** thinks tx<sub>1</sub> is invalid.

# We need a better protocol!

What is the form of such protocol?

We need to define *what code each node will run*

The code must:

- **Injection:** Allow “publishing” a local *tx*
- **Reading:** When locally asked “what tx have happened?”, returns a **ledger**

*A ledger is an ordered list of tx*

This is the **local view** of the honest party

# Desired properties of a ledger

## 1. Liveness

If I “inject” a transaction, it will *eventually* appear in the *ledger* of *all honest parties*

## 2. Persistence

If a transaction appears in a certain position  $i$  within an honest party's ledger, it will appear *in the same position* for *all* honest parties

# Liveness & Persistence

- The solution is easy if we give up one of the properties. **How?**

- Without **liveness**:

**Always return an empty ledger.**

tx injected will *never* appear in any ledger.

Persistence is trivially satisfied.

- Without **persistence**:

**Append all tx you see on network to ledger immediately.**

Parties won't agree about order of tx.

Liveness is trivially satisfied.


# Assumption: Total connectivity

Each node is **connected to everyone** else

In particular: **all honest parties are directly connected** to each other

Not true in practice, but realistic model:

All nodes are connected *indirectly*; a message is broadcast, then relayed. Allows us to abstract out the notion of relaying.



# Summary: The network model

1. The network is **synchronous**:

A message by an honest party at round  $r$  is delivered to all parties at  $r + 1$

2. The network is **reliable**:

The adversary cannot drop honest messages

3. The network is **Sybil attackable**:

The adversary can inject arbitrarily many messages

The adversary can send different messages to different parties

4. Parties **do not have pre-known identities**:


The adversary can reorder and shuffle network messages



# The execution

- Every party is modelled as an Interactive Turing Machine
- Execution is coordinated by a special PPT TM, the *environment Z*
- The environment captures the notion of “rounds” (an iteration of  $Z$ )
- For every round, it invokes *each honest party*
  - It gives it network messages from the previous round
  - It receives from it a network message to be delivered in the next round
- At the end of each round, it invokes the adversary (a “*rushing adversary*”)
- The adversary can see all honest party messages, reorder them, or inject fake ones (Sybil attack), but not delete them. The changes may be different for each honest party.
- The new modified messages are delivered to honest parties at the next round
- Parties can keep local state

## Environment Z




# The environment

```
tampered_messages = [ $\varepsilon$ ,  $\varepsilon$ , ...,  $\varepsilon$ ]  
for round = 1 to poly( $\kappa$ )  
  for i = 1 to n - t  
    next_messages[i] = P[i](tampered_messages[i])  
  end for  
  tampered_messages = A(next_messages)  
  for m in next_messages  
    for i = 1 to n - t  
      assert(m in tampered_messages[i])  
    end for  
  end for  
end for
```

# The random oracle (RO) model

- A truly random function  $H$  with output  $\{0, 1\}^k$  is chosen before the beginning of the execution
- Any party can query the function
- The function is shared among all honest parties and the adversary
- It gives the *same* output for the same input even if asked by different parties
- We think of it as a *cryptographically secure hash function*



# The honest majority assumption

- We try to capture that the adversary has *less computational power* than the honest parties combined
- This is how we will distinguish the honest opinion from the adversarial opinion
- Computational power cannot be Sybil attacked
- We allow  $q$  RO queries per round for each honest party
- We allow the adversary  $tq$  RO queries per round
- We require:

adversarial parties

honest parties

$$t \leq (1 - \delta)(n - t)$$

a small gap

(the honest majority equation)

# Summary: Our assumptions

- The network is **synchronous, connected, Sybil-attackable**
- The analysis is in the **random oracle model**
- The **majority** of computational power is **honest at all times**
- **We don't analyze incentives**
- **The number of parties is constant**

Many of these assumptions can be lifted, but the analysis becomes more complex (e.g. the “constant number of parties” assumption is relaxed in follow-up work)

# Proof-of-Work (PoW)

How can we prove that we hold certain computational power?

- Give a *fresh input* to RO
- If the output is sufficiently small, we have (probabilistically) proven that we have expended some of our  $q$  RO queries
- We can use this mechanism for *voting*
- Let  $T$  be the bound required, then we want:


$$H(x) < T$$

(the proof-of-work equation)

- $H(x)$  and  $T$  are both  $\kappa$ -bit strings compared as binary numbers

# A block


- Contains an ordered sequence of transactions
- The order must be valid (no double spends)
- Contains some proof-of-work, proving computational power has been expended in creating it:  $H(\mathbf{B}) < T$






# A blockchain

- A sequence of blocks
- Each block contains the hash of its parent
- **A block could not have been generated before its parent (Why?)**
- Therefore, blocks are generated in order
- A block buried under sufficient other blocks is **voted on** by computational power





# Mining

- Each honest party *adopts* a blockchain  $C$
- Not all honest parties necessarily adopt the same chain  
(we would like that, but we cannot ensure it due to the consensus problem)
- During each round, each honest party attempts to extend their adopted chain
- ...by creating a new block  $B$  on top of  $C$  such that  $CB$  is a valid chain
- $B$  points to  $C[-1]$
- $B$  must contain valid PoW. It contains some nonce  $ctr$  to achieve the PoW.
- $B$  contains a vector of tx  $x$  that the honest party wishes to inject and have not yet been included
- If honest party is successful in finding  $B$ , it broadcasts  $CB$  to the network

**Algorithm 3** The *proof of work* function, parameterized by  $q$ ,  $T$  and hash functions  $H(\cdot), G(\cdot)$ . The input is  $(x, \mathcal{C})$ .

---

```
1: function pow( $(x, \mathcal{C})$ )
2:   if  $\mathcal{C} = \varepsilon$  then injection ▷ Determine proof of work instance
3:      $s \leftarrow 0$ 
4:   else
5:      $\langle s', x', ctr' \rangle \leftarrow \text{head}(\mathcal{C})$ 
6:      $s \leftarrow H(ctr', G(s', x'))$ 
7:   end if
8:    $ctr \leftarrow 1$ 
9:    $B \leftarrow \varepsilon$ 
10:   $h \leftarrow G(s, x)$ 
11:  while ( $ctr \leq q$ ) do
12:    if ( $H(ctr, h) < T$ ) then ▷ This  $H(\cdot)$  invocation subject to the  $q$ -bound
13:       $B \leftarrow \langle s, x, ctr \rangle$  PoW equation
14:      break
15:    end if
16:     $ctr \leftarrow ctr + 1$ 
17:  end while
18:   $\mathcal{C} \leftarrow \mathcal{C}B$  ▷ Extend chain
19:  return  $\mathcal{C}$ 
20: end function
```

---

# Validating a chain

- Check that the transactions it contains form a valid sequence
- Check the proof-of-work of each block
- Check that each block points to its parent

---

**Algorithm 1** The *chain validation predicate*, parameterized by  $q, T$ , the hash functions  $G(\cdot), H(\cdot)$ , and the *content validation predicate*  $V(\cdot)$ . The input is  $\mathcal{C}$ .

---

```
1: function validate( $\mathcal{C}$ )
2:    $b \leftarrow V(\mathbf{x}_{\mathcal{C}})$  validate transaction sequence
3:   if  $b \wedge (\mathcal{C} \neq \varepsilon)$  then ▷ The chain is non-empty and meaningful w.r.t.  $V(\cdot)$ 
4:      $\langle s, x, ctr \rangle \leftarrow \text{head}(\mathcal{C})$ 
5:      $s' \leftarrow H(ctr, G(s, x))$ 
6:     repeat
7:        $\langle s, x, ctr \rangle \leftarrow \text{head}(\mathcal{C})$ 
8:       if  $\text{validblock}_q^T(\langle s, x, ctr \rangle) \wedge (H(ctr, G(s, x)) = s')$  then validate chain ancestry
9:          $s' \leftarrow s$  ▷ Retain hash value
10:         $\mathcal{C} \leftarrow \mathcal{C}^{\lceil 1}$  validate block PoW ▷ Remove the head from  $\mathcal{C}$ 
11:      else
12:         $b \leftarrow \text{False}$ 
13:      end if
14:    until  $(\mathcal{C} = \varepsilon) \vee (b = \text{False})$ 
15:  end if
16:  return ( $b$ )
17: end function
```

---

# The Longest Chain Rule

- Which chain should I adopt?
- The longest chain! (In terms of **number of blocks**, not number of tx!)
- Why? It contains the most accumulated proof-of-work
- This signifies that a lot of computational power is vouching for it
- Hence it will be what the honest parties believe!  
(intuitively -- exact property formulated below)

---

**Algorithm 2** The function that finds the “best” chain, parameterized by function  $\max(\cdot)$ . The input is  $\{\mathcal{C}_1, \dots, \mathcal{C}_k\}$ .

---

```
1: function maxvalid( $\mathcal{C}_1, \dots, \mathcal{C}_k$ )
2:    $temp \leftarrow \varepsilon$ 
3:   for  $i = 1$  to  $k$  do
4:     if validate( $\mathcal{C}_i$ ) then
5:        $temp \leftarrow \max(\mathcal{C}_i, temp)$ 
6:     end if
7:   end for
8:   return  $temp$ 
9: end function
```

---



# The blockchain backbone protocol

- Maintain a blockchain
- When seeing other blockchains on the network, adopt the longest one
- Mine on top of longest chain
- If mining successful, broadcast mined chain

**Algorithm 4** The Bitcoin backbone protocol, parameterized by the *input contribution function*  $I(\cdot)$  and the *chain reading function*  $R(\cdot)$ . At the onset it is assumed “init= True”.


---

```
1: if (init) then
2:    $\mathcal{C} \leftarrow \varepsilon$ 
3:    $st \leftarrow \varepsilon$ 
4:    $round \leftarrow 1$ 
5:    $init \leftarrow \text{False}$ 
6: else
7:    $\tilde{\mathcal{C}} \leftarrow \text{maxvalid}(\mathcal{C}, \text{any chain } \mathcal{C}' \text{ found in RECEIVE}())$ 
8:   if INPUT() contains READ then
9:     write  $R(\tilde{\mathcal{C}})$  to OUTPUT()  $\triangleright$  Produce necessary output before the POW stage.
10:  end if ledger reading
11:   $\langle st, x \rangle \leftarrow I(st, \tilde{\mathcal{C}}, round, \text{INPUT}(), \text{RECEIVE}())$   $\triangleright$  Determine the  $x$ -value.
12:   $\mathcal{C}_{\text{new}} \leftarrow \text{pow}(x, \tilde{\mathcal{C}})$  injection
13:  if  $\mathcal{C} \neq \mathcal{C}_{\text{new}}$  then
14:     $\mathcal{C} \leftarrow \mathcal{C}_{\text{new}}$ 
15:    DIFFUSE( $\mathcal{C}$ )  $\triangleright$  Broadcast the chain in case of adoption/extension.
16:  else
17:    DIFFUSE( $\perp$ )  $\triangleright$  Signals the end of the round to the diffuse functionality.
18:  end if
19:   $round \leftarrow round + 1$ 
20: end if
```

---


# An honest execution

- Consider: What happens if all parties are honest?
- An honest party finds a block, broadcasts it
- The rest of the network adopts the new chain, as it is longer
- Another honest party finds a new block, broadcasts it...
- The chain keeps growing and is the same among all honest parties
- The rate of growth corresponds to the parameters  $T$ ,  $q$ ,  $n$



# Blockchain forks

- Sometimes two blocks will be mined at the same round by two honest parties
- Any of the two alternatives can be adopted by honest parties
  - Which one is chosen by the adversary, as they control the network!
- The tie will be broken when a next block is found
- Then the longest chain will be unique again
- This is why we want the probability of block finding to be small



# Successful rounds

We call a round...

- *successful* if (at least) an **honest** party finds a block
- *uniquely successful* if **exactly one** honest party finds a block

It's possible that the adversary will also find blocks during successful or uniquely successful rounds.

$$f \stackrel{\text{def}}{=} \Pr[ \text{round } r \text{ is successful } ]$$

$$= 1 - (1 - T / 2^K)^{q(n - t)}$$

$$\cong q(n - t) T / 2^K$$

# Blockchains have three security properties

1. **Chain growth**

The chain adopted by an honest party grows

2. **Chain quality**

Large chunks of an honestly adopted chain contain honest blocks

3. **Common prefix**

All honestly adopted chains share a large common prefix

Theorem: **For all PPT adversaries**  $A$ , the properties hold with *overwhelming probability* in the security parameter.

# Chain growth

Consider an honest party P.

If P has adopted chain  $C_1$  at round  $r$ , and  $C_2$  at round  $r + s$ , then

$$|C_2| \geq |C_1| + \tau s$$

(the chain growth equation)

$\tau$  is the **chain velocity**. This property holds for sufficiently large  $s$ .



# Intuition for Chain Growth

Consider a successful round with a block by party  $P_1$ .

During a successful round, the chain of all honest parties grows.

- The chain of  $P_1$  grows by the new block
- Consider  $P_2$ 
  - If  $P_2$  is successful, their chain also grows by *their* new block
  - If  $P_2$  is unsuccessful, their chain grows because they receive  $P_1$ 's extension
- The adversary cannot harm this due to Longest Chain Rule

# Intuition for Chain Growth

- Consider a sequence of  $s$  rounds
- Each round has a probability  $f$  of being successful
- In expectation,  $f s$  rounds will be successful
- Hence the chain will grow by approximately  $f s$  blocks
- Therefore  $\tau \approx f$

The full proof is by induction on the number of rounds  
and a Chernoff bound with overwhelming probability in  $s$

# Chain quality

Consider an honestly adopted chain  $C$ .

For any  $i$ ,

**$C[i: i + \ell]$  contains at least  $\mu \ell$  honest blocks**

$\mu$  is the proportion of honest blocks. This property holds for sufficiently large  $\ell$ .

# Common prefix

Let  $P_1, P_2$  be two honest parties and consider a round  $r$  at which their adopted chains are  $C_1$  and  $C_2$ . Then  $C_1[: -k]$  is a prefix of  $C_2$  and vice versa:


$$C_2[ : |C_1[: -k]| ] = C_1[: -k]$$


(the common prefix equation)


This property holds for sufficiently large  $k$ .

# Intuition for Common Prefix

- Consider an adversary who double spends
- They have to extend a chain faster than the honest parties
- This is impossible if the honest chain is large enough







# Implementing the Ledger functionality

Now that we have a blockchain, it's easy...

- Take `C[:-k]` and obtain its transactions
- Return that ledger!



# Proof of Liveness

Suppose we wish to inject a *tx* starting at round  $r$

- By the *Chain Growth* property, the chain will keep growing
- After  $s$  rounds, there will be a growth by  $\tau s$  new blocks
- For sufficiently large  $s$  (“eventually”), we have that  $\tau s \mu \geq 1$
- Therefore, after  $s$  rounds, an honestly generated block  $B$  will be adopted by all honest parties
- Due to *Chain Growth*,  $B$  will eventually be buried under  $k$  blocks
- It will then be reported in the ledger

# Proof of Persistence

Suppose a transaction is reported in the ledger of an honest party  $P_1$


- Then it is in  $C_1[:-k]$  of  $P_1$
- Suppose it is reported at a different position by some other party  $P_2$
- Then it is in  $C_2[:-k]$  of  $P_2$
- But this means  $C_1[:-k]$  does not share a prefix with  $C_2[:-k]$
- This violates the Common Prefix property!

# References

1. Juan Garay, Aggelos Kiayias, and Nikos Leonardos.  
[The bitcoin backbone protocol: Analysis and applications](#).  
In Annual International Conference on the Theory and Applications of Cryptographic Techniques, pages 281–310. Springer, 2015.
2. Juan Garay, Aggelos Kiayias, and Nikos Leonardos.  
*The bitcoin backbone protocol with chains of variable difficulty*.  
In Annual International Cryptology Conference, pages 291–323. Springer, 2017.

45DC 00AE FDDF 5D5C B988 EC86 2DA4 50F3 AFB0 46C7

# Thanks! Questions?



INPUT | OUTPUT



---

National and Kapodistrian  
UNIVERSITY OF ATHENS

---



Only some rights reserved